



Text Data Vectorization – From Text to Numbers

SOFIA
UNIVERSITY



ST. KLIMENT
OHRIDSKI
1888


FACULTY
OF ECONOMICS
AND BUSINESS
ADMINISTRATION

What you will learn

- ◎ What is **Text Vectorization**? 
- ◎ Why do we need it? 
- ◎ **Key techniques** for text data vectorization. 
- ◎ How to **apply** them in **Python**? 
- ◎ More advanced techniques for text data vectorization. 

1.

From Text... to Numbers



Text Vectorization (1)

◎ Formal definition of **text vectorization**:

“


*The process of converting textual data into numerical form is called ‘**text vectorization**’.*

In this way, words, sentences or whole documents in a corpus are represented as vectors of numerical values.

”



Text Vectorization (2)



Vocabulary = [Word 1, Word 2, Word 3]

Text Vectorization (3)

- ◎ Quantitative analysis **can't** be applied without text vectorization!
- ◎ What **form of text vectorization** to apply depends on the problem at hand and the chosen algorithm for data analysis.
- ◎ Text vectorization might **highly impact** the results of the analysis!
- ◎ There are **basic and more complex** methods for text vectorization.

The Vector Space Model (1)

In the vector space model each document \mathbf{d}_i is represented as a vector of weights \mathbf{u}_{ij} :

$$\mathbf{d}_i = (\mathbf{u}_{i1}, \mathbf{u}_{i2} \dots, \mathbf{u}_{iV}), \text{ where}$$

\mathbf{D} – the sample of text data.

M – the number of documents in sample \mathbf{D} .

\mathbf{d}_i – a given document in \mathbf{D} , $i = 1 \dots M$.

\mathbf{w}_j – a given token part of the vocabulary, $j = 1 \dots V$.

V – number of unique tokens in the sample. The unique tokens form the “**vocabulary**” of the model.

\mathbf{u}_{ij} – the weight of the j^{th} token \mathbf{w}_j in the document \mathbf{d}_i , where $(i = 1 \dots M)$ and $(j = 1 \dots V)$.

The Vector Space Model (2)

Input Text:
“the plot was
good”



Vocabulary:

the
plot
was
very
good
boring



**Vector
Representation:**

$$\begin{bmatrix} 1 \\ 1 \\ 1 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

The Bag-of-Words assumption

- ◎ The **bag-of-words assumption** conveys the idea of a text representation that disregards the order of words.
- ◎ The vector-space model makes the bag-of-words assumption - **order of words in text documents is not taken into account.**
- ◎ According to this assumption “Mary is quicker than John” **is the same as** “John is quicker than Mary” (Manning, 2008).
- ◎ Despite the “unrealistic” nature of this assumption, it proves to **work well in many practical text mining tasks.**
- ◎ Making this assumption we say: “**word order is not important but I expect that documents with similar representation share similar content**”.

Forms of Vectorization in the Vector Space Model

There are **three forms of text vectorization** in the Vector Space Model:

- ◎ Binary
- ◎ Count (absolute or relative)
- ◎ TF-IDF (term frequency – inverse document frequency)

Binary Vectorization

- Binary vectorization takes into account only the absence or presence of a given token. The **presence** is denoted with **1**, while the **absence** is denoted with **0**.

“the movie is ok”

“the beginning was boring...very boring”

“the plot was very good”

Our vocabulary



	“beginning”	“boring”	“good”	“is”	“movie”	“ok”	“plot”	“the”	“very”	“was”
Review 1	0	0	0	1	1	1	0	1	0	0
Review 2	1	1	0	0	0	0	0	1	1	1
Review 3	0	0	1	0	0	0	1	1	1	1

Count Vectorization (absolute)

- Count vectorization takes into account the **frequency of tokens** in the documents.
- Example of count vectorization with **absolute values**:

“the movie is ok”

“the beginning was boring...very boring”

“the plot was very good”

	“beginning”	“boring”	“good”	“is”	“movie”	“ok”	“plot”	“the”	“very”	“was”
Review 1	0	0	0	1	1	1	0	1	0	0
Review 2	1	2	0	0	0	0	0	1	1	1
Review 3	0	0	1	0	0	0	1	1	1	1

Count Vectorization (relative)

- ◎ In relative count vectorization, the **counts are normalized** by the number of words in the document.
- ◎ Example of relative count vectorization:

“the movie is ok”

“the beginning was boring...very boring”

“the plot was very good”

	“beginning”	“boring”	“good”	“is”	“movie”	“ok”	“plot”	“the”	“very”	“was”
Review 1	0	0	0	0.25	0.25	0.25	0	0.25	0	0
Review 2	0.16	0.33	0	0	0	0	0	0.16	0.16	0.16
Review 3	0	0	0.2	0	0	0	0.2	0.2	0.2	0.2

Keywords are important, right?

Mayday! And May Day! (*New York Times*, 1st of May 1986)

Those who live by **secrecy** can also perish by it. The **Chernobyl** nuclear **disaster** may have begun as early as **last Friday**, but the **Soviet Union** suppressed all **news** of it until **Sweden** reported **radiation** on **Monday**. That **delay** in warning neighboring **country** of the **impending catastrophe** alarmed and misled **people** from the Elbe to the Urals. **Mikhail Gorbachev** cannot win **confidence** in his **pledge** to reduce nuclear **weapon** if he forfeits his **neighbor** **trust** over the **peaceful use** of nuclear **energy**.

Topics:

- 3.Cold War
- 8.Accidents
- 1.Research
- 7.Proliferation
- 9.Nuclear power

...TF-IDF vectorization?

TF-IDF Vectorization

Term frequency

Inverse-document frequency

$$TF-IDF(w_n, d_i) = tf_{w_n, d_i} \times idf_{w_n}, \text{ where:}$$

D – the sample of text data.

M – the number of documents in sample D .

d_i – a given document in D , $i = 1 \dots M$.

w_n – n^{th} word in document d_i , $n \in \{1, \dots, N_d\}$.

$$df_{w_n} = |\{d_i \in D : w_n \in d_i\}|$$

$$idf_{w_n} = \log \frac{M}{df_{w_n}}$$

$$TF-IDF(w_n, d_i) = tf_{w_n, d_i} \times \log \frac{M}{df_{w_n}}$$

TF-IDF Vectorization in scikit-learn

- © The implementation in scikit-learn uses the following formula for calculating TF-IDF:

tf_{w_n, d_i} - number of times the word occurs in the given document

df_{w_n} - number of documents in the corpus containing the word

$$idf_{w_n} = \ln \frac{1 + M}{1 + df_{w_n}} + 1$$

$$\mathbf{TF-IDF}(w_n, d_i) = tf_{w_n, d_i} \times idf_{w_n}$$

NB! L2 normalization is applied on the vectors of TF-IDF values (sum of squares = 1).



Original vector: $x = [x_1 \ x_2 \ \dots \ x_V]$

$$y_1 = \frac{x_1}{\sqrt{\sum_{i=1}^V x_i^2}}$$

L2 normalized vector: $y = [y_1 \ y_2 \ \dots \ y_V], y_i \in [0,1]$

TF-IDF Calculation (1)

“the movie is ok”
“the beginning was boring...very boring”
“the plot was very good”

Calculate TF-IDF values for the first review:

$$\text{“the”} \Rightarrow TF = 1, DF = 3 \Rightarrow idf = \ln \frac{1+3}{1+1} + 1 = 1 \Rightarrow tf \times idf = 1 \times 1 = 1$$

$$\text{“movie”} \Rightarrow TF = 1, DF = 1 \Rightarrow idf = \ln \frac{1+3}{1+1} + 1 = 1.693 \Rightarrow tf \times idf = 1 \times 1.693 = 1.693$$

$$\text{“is”} \Rightarrow TF = 1, DF = 1 \Rightarrow idf = \ln \frac{1+3}{1+1} + 1 = 1.693 \Rightarrow tf \times idf = 1 \times 1.693 = 1.693$$

$$\text{“ok”} \Rightarrow TF = 1, DF = 1 \Rightarrow idf = \ln \frac{1+3}{1+1} + 1 = 1.693 \Rightarrow tf \times idf = 1 \times 1.693 = 1.693$$

First review before normalization:

[0 0 0 1.693 1.693 1.693 0 1 0 0]

TF-IDF Calculation (2)

“the movie is ok”
“the beginning was boring...very boring”
“the plot was very good”

By default in scikit-learn L2 normalization is applied on the TF-IDF weights (recommended approach):

First review:

$$\text{“the”} \Rightarrow \text{TF-IDF} = 1 \Rightarrow \frac{1}{\sqrt{1^2 + 1.693^2 + 1.693^2 + 1.693^2}} = \frac{1}{3.098} \sim 0.3227$$

$$\text{“movie”} \Rightarrow \text{TF-IDF} = 1.693 \Rightarrow \frac{1.693}{\sqrt{1^2 + 1.693^2 + 1.693^2 + 1.693^2}} = \frac{1.693}{3.098} \sim 0.5464$$

$$\text{“is”} \Rightarrow \text{TF-IDF} = 1.693 \Rightarrow \frac{1.693}{\sqrt{1^2 + 1.693^2 + 1.693^2 + 1.693^2}} = \frac{1.693}{3.098} \sim 0.5464$$

$$\text{“ok”} \Rightarrow \text{TF-IDF} = 1.693 \Rightarrow \frac{1.693}{\sqrt{1^2 + 1.693^2 + 1.693^2 + 1.693^2}} = \frac{1.693}{3.098} \sim 0.5464$$

First review after normalization:

[0 0 0 0.5464 0.5464 0.5464 0 0.3227 0 0]

TF-IDF Calculation (3)

“the movie is ok”
“the beginning was boring...very boring”
“the plot was very good”

Final matrix of TF-IDF values calculated in scikit-learn:

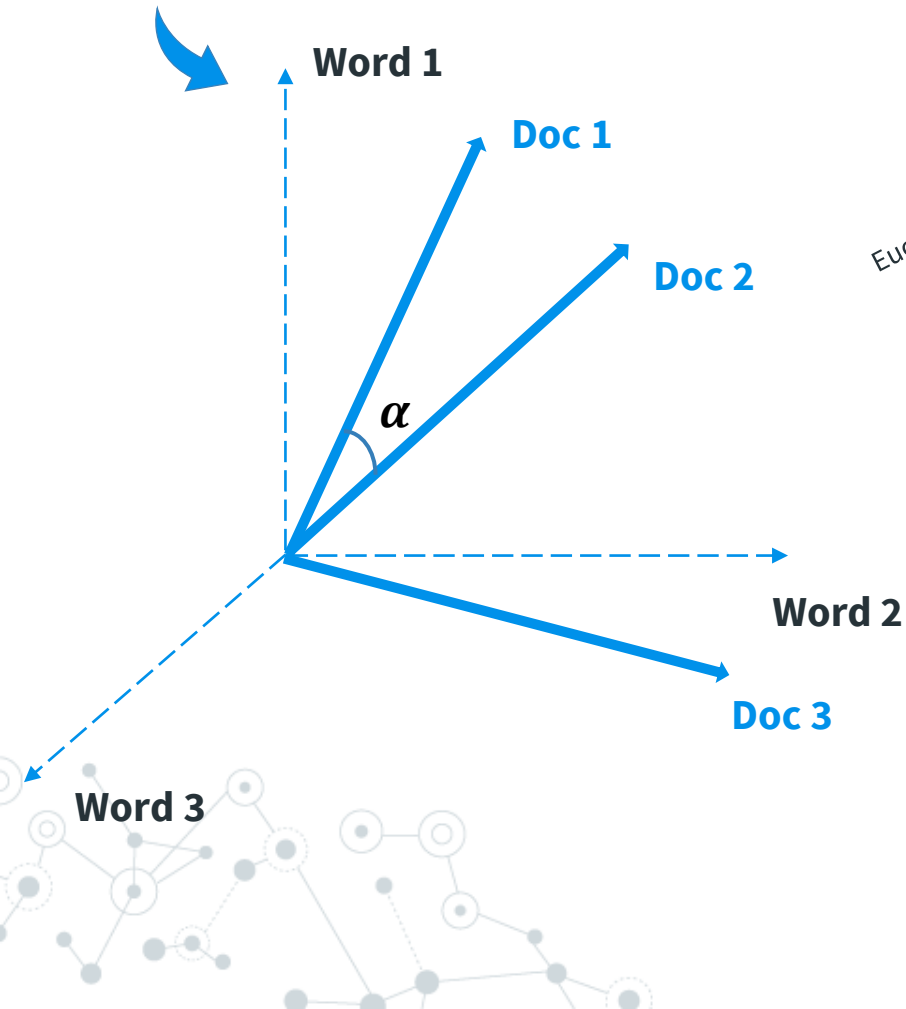
	“beginning”	“boring”	“good”	“is”	“movie”	“ok”	“plot”	“the”	“very”	“was”
Review 1	0	0	0	0.5464	0.5464	0.5464	0	0.3227	0	0
Review 2	0.3920	0.7841	0	0	0	0	0	0.2315	0.2981	0.2981
Review 3	0	0	0.5340	0	0	0	0.5340	0.3154	0.4061	0.4061

Practical Applications of the Vector Space Model

- ◎ **Text classification/clustering** (sentiment analysis, topic modeling etc.);
- ◎ Fundamental in **Information Retrieval** – serves as a basis for search engines development;
- ◎ **Recommender systems** – for movies, books, news articles etc.
- ◎ And many other.

Find Text Similarity with TF-IDF

Vocabulary = [Word 1, Word 2, Word 3]



Cosine similarity:

$$\text{sim}(d_1, d_2) = \frac{\vec{V}(d_1) \cdot \vec{V}(d_2)}{|\vec{V}(d_1)| |\vec{V}(d_2)|}$$

Euclidean vector length

$$|\vec{V}(d_1)| = \sqrt{\sum_{i=1}^V \vec{V}_i^2(d_1)}$$

$$\frac{\vec{V}(d_1)}{|\vec{V}(d_1)|} = \frac{\vec{V}(d_1)}{\sqrt{\sum_{i=1}^V \vec{V}_i^2(d_1)}} = \vec{v}(d_1)$$

$$\text{sim}(d_1, d_2) = \vec{v}(d_1) \cdot \vec{v}(d_2)$$

$$\text{sim}(d_1, d_2) = \cos \alpha$$




How to implement all these techniques in Python?

Use scikit-learn (<https://scikit-learn.org/stable/index.html>):


- ◎ Binary vectorization –
CountVectorizer(binary= True)
- ◎ Count Vectorization - CountVectorizer()
- ◎ TF-IDF Vectorization – TfidfVectorizer()
- ◎ Cosine similarity - cosine_similarity()

More complex methods for text vectorization


- Word Embeddings – word2vec (2013), doc2vec, GloVe and other.
- An upgrade in 2018 - ELMo and BERT



Male-Female



Verb tense



Country-Capital

*„You shall know a word by the company it keeps...”
(Firth, J. R. 1957:11)*

If you want to learn more... 🤖

- ◎ [Introduction to Information Retrieval \(stanford.edu\)](#) – p.154 - p.161 (pdf)
- ◎ Miner et al., “**Practical Text Mining and Statistical Analysis for Non-structured Text Data Applications**” - Chapter 3



Thanks!

Any questions?

You can find me at:
`g.hristova@feb.uni-sofia.bg`

